

# LasVegas

Andie

9/1/2017

```
#install.packages('corrplot')
#install.packages('doBy')
library(doBy) #summaryBy
library(ggplot2) # Data visualization
library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(corrplot) # correlation plot
library(plotly) # interactive map

##
## Attaching package: 'plotly'

## The following object is masked from 'package:ggplot2':
##
##   last_plot

## The following object is masked from 'package:stats':
##
##   filter

## The following object is masked from 'package:graphics':
##
##   layout

library(rworldmap) #world map

## Loading required package: sp

## ### Welcome to rworldmap ###

## For a short introduction type :   vignette('rworldmap')
```

Lets open our file. We notice that its a csv file and it uses semi colons to separate the values, so we will use read.csv2 to open our file and assign it to 'the 'df'.

```
df <- read.csv2('/Users/andiedonovan/myProjects/LasVegas/LasVegas.csv') # open file
#View(df) # view the whole data set
```

We notice that some of the column names are long or messy, so lets fix that:

There are a lot of columns, so it takes a long time to rename the columns. Say we didn't really care about their length, but just wanted to remove the periods from the column names. We could use the function `gsub` to easily do this:

```
names(df) <- gsub("\\.", "", names(df))
head(df)
```

	Usercountry	Nrreviews	Nrhotelreviews	Helpfulvotes	Score	Periodofstay
## 1	USA	11	4	13	5	Dec-Feb
## 2	USA	119	21	75	3	Dec-Feb
## 3	USA	36	9	25	5	Mar-May
## 4	UK	14	7	14	4	Mar-May
## 5	Canada	5	5	2	4	Mar-May
## 6	Canada	31	8	27	3	Mar-May

	Travelertype	Pool	Gym	Tenniscourt	Spa	Casino	Freeinternet
## 1	Friends	NO	YES	NO	NO	YES	YES
## 2	Business	NO	YES	NO	NO	YES	YES
## 3	Families	NO	YES	NO	NO	YES	YES
## 4	Friends	NO	YES	NO	NO	YES	YES
## 5	Solo	NO	YES	NO	NO	YES	YES
## 6	Couples	NO	YES	NO	NO	YES	YES

		Hotelname	Hotelstars	Nrrooms	Usercontinent
## 1	Circus	Circus Hotel & Casino Las Vegas	3	3773	North America
## 2	Circus	Circus Hotel & Casino Las Vegas	3	3773	North America
## 3	Circus	Circus Hotel & Casino Las Vegas	3	3773	North America
## 4	Circus	Circus Hotel & Casino Las Vegas	3	3773	Europe
## 5	Circus	Circus Hotel & Casino Las Vegas	3	3773	North America
## 6	Circus	Circus Hotel & Casino Las Vegas	3	3773	North America

	Memberyears	Reviewmonth	Reviewweekday
## 1	9	January	Thursday
## 2	3	January	Friday
## 3	2	February	Saturday
## 4	6	February	Friday
## 5	7	March	Tuesday
## 6	2	March	Tuesday

If we did want to go ahead and rename all of the columns, we could do that using `names`, `colnames` and a list:

```
#Using names
names(df) # our current column names
```

## [1]	"Usercountry"	"Nrreviews"	"Nrhotelreviews"	"Helpfulvotes"
## [5]	"Score"	"Periodofstay"	"Travelertype"	"Pool"
## [9]	"Gym"	"Tenniscourt"	"Spa"	"Casino"
## [13]	"Freeinternet"	"Hotelname"	"Hotelstars"	"Nrrooms"
## [17]	"Usercontinent"	"Memberyears"	"Reviewmonth"	"Reviewweekday"

```
names(df)[1]<-"Country" #rename just first one column
head(df) #check that the first column name changed from 'Usercountry' to 'Country'
```

	Country	Nrreviews	Nrhotelreviews	Helpfulvotes	Score	Periodofstay
## 1	USA	11	4	13	5	Dec-Feb
## 2	USA	119	21	75	3	Dec-Feb
## 3	USA	36	9	25	5	Mar-May
## 4	UK	14	7	14	4	Mar-May
## 5	Canada	5	5	2	4	Mar-May

```
## 6 Canada 31 8 27 3 Mar-May
## Travelertype Pool Gym Tenniscourt Spa Casino Freeinternet
## 1 Friends NO YES NO NO YES YES
## 2 Business NO YES NO NO YES YES
## 3 Families NO YES NO NO YES YES
## 4 Friends NO YES NO NO YES YES
## 5 Solo NO YES NO NO YES YES
## 6 Couples NO YES NO NO YES YES
## Hotelname Hotelstars Nrrooms Usercontinent
## 1 Circus Circus Hotel & Casino Las Vegas 3 3773 North America
## 2 Circus Circus Hotel & Casino Las Vegas 3 3773 North America
## 3 Circus Circus Hotel & Casino Las Vegas 3 3773 North America
## 4 Circus Circus Hotel & Casino Las Vegas 3 3773 Europe
## 5 Circus Circus Hotel & Casino Las Vegas 3 3773 North America
## 6 Circus Circus Hotel & Casino Las Vegas 3 3773 North America
## Memberyears Reviewmonth Reviewweekday
## 1 9 January Thursday
## 2 3 January Friday
## 3 2 February Saturday
## 4 6 February Friday
## 5 7 March Tuesday
## 6 2 March Tuesday
```

*#Using colnames*

`colnames(df)` *# gives us the same thing*

```
## [1] "Country" "Nrreviews" "Nrhotelreviews" "Helpfulvotes"
## [5] "Score" "Periodofstay" "Travelertype" "Pool"
## [9] "Gym" "Tenniscourt" "Spa" "Casino"
## [13] "Freeinternet" "Hotelname" "Hotelstars" "Nrrooms"
## [17] "Usercontinent" "Memberyears" "Reviewmonth" "Reviewweekday"
```

`colnames(df)<-c('Country', 'NoReviews', 'NoHotelReviews', 'Helpful', 'Score', 'Stay', 'Traveler', 'Pool', 'Gym', 'Tennis', 'Spa', 'Casino', 'Internet', 'Name', 'Stars', 'NoRms', 'Continent', 'MemberYrs', 'Month', 'Weekday')` *#rename all of the columns using a list*

`head(df)` *# check that it worked!*

```
## Country NoReviews NoHotelReviews Helpful Score Stay Traveler Pool Gym
## 1 USA 11 4 13 5 Dec-Feb Friends NO YES
## 2 USA 119 21 75 3 Dec-Feb Business NO YES
## 3 USA 36 9 25 5 Mar-May Families NO YES
## 4 UK 14 7 14 4 Mar-May Friends NO YES
## 5 Canada 5 5 2 4 Mar-May Solo NO YES
## 6 Canada 31 8 27 3 Mar-May Couples NO YES
## Tennis Spa Casino Internet Name Stars
## 1 NO NO YES YES Circus Circus Hotel & Casino Las Vegas 3
## 2 NO NO YES YES Circus Circus Hotel & Casino Las Vegas 3
## 3 NO NO YES YES Circus Circus Hotel & Casino Las Vegas 3
## 4 NO NO YES YES Circus Circus Hotel & Casino Las Vegas 3
## 5 NO NO YES YES Circus Circus Hotel & Casino Las Vegas 3
## 6 NO NO YES YES Circus Circus Hotel & Casino Las Vegas 3
## NoRms Continent MemberYrs Month Weekday
## 1 3773 North America 9 January Thursday
## 2 3773 North America 3 January Friday
```

```
## 3 3773 North America      2 February Saturday
## 4 3773      Europe        6 February  Friday
## 5 3773 North America      7   March   Tuesday
## 6 3773 North America      2   March   Tuesday
```

There's a lot of variables in our dataset, and we probably will not need all of them for our analysis. We can create different sub datasets to allow for easier analysis of specific factors. For example, to see how whether or not a hotel has certain amenities affects its score, let's make a dataset that only includes the hotel name, country, score, and the 6 amenities listed. Let's call this new dataset `amenities1`:

Also notice the values for each of the variables is either 'YES' or 'NO'. It's much easier to work with numbers than characters, so let's turn each of the variables into a binary factor with 'YES'=2 and 'NO'=1

```
amenities1 = df %>%
  select(Name, Country, Score, Pool, Gym, Tennis, Spa, Casino, Internet)

amenities1$Pool<-as.numeric(amenities1$Pool)
amenities1$Gym<-as.numeric(amenities1$Gym)
amenities1$Tennis<-as.numeric(amenities1$Tennis)
amenities1$Spa<-as.numeric(amenities1$Spa)
amenities1$Casino<-as.numeric(amenities1$Casino)
amenities1$Internet<-as.numeric(amenities1$Internet)

head(amenities1)
```

	Name	Country	Score	Pool	Gym	Tennis	Spa
## 1	Circus Circus Hotel & Casino Las Vegas	USA	5	1	2	1	1
## 2	Circus Circus Hotel & Casino Las Vegas	USA	3	1	2	1	1
## 3	Circus Circus Hotel & Casino Las Vegas	USA	5	1	2	1	1
## 4	Circus Circus Hotel & Casino Las Vegas	UK	4	1	2	1	1
## 5	Circus Circus Hotel & Casino Las Vegas	Canada	4	1	2	1	1
## 6	Circus Circus Hotel & Casino Las Vegas	Canada	3	1	2	1	1
##	Casino	Internet					
## 1	2	2					
## 2	2	2					
## 3	2	2					
## 4	2	2					
## 5	2	2					
## 6	2	2					

Above you'll see a symbol that looks like this: `%>%`. This is called a pipe operator (dplyr package) and is used to insert some input or argument into a function. It's useful when we have sequences of operations and can help make the flow of the steps more readable/ easier to follow.

Equivalently, we could nest the above code and get the same results. This would look like: `amenities = select(df, c(Name, Country, Score, Pool, Gym, Tennis, Spa, Casino, Internet))`. Although in this case, it doesn't make much of a difference which format you use, later on, the pipe operator is a powerful tool for applying multiple operations to one object.

Alternative data set without hotel name & country (ie only numerical):

```
amenities2 = df %>%
  select(Score, Pool, Gym, Tennis, Spa, Casino, Internet) # columns to be selected

amenities2 = amenities2 %>% mutate(Pool=as.factor(ifelse(Pool=="YES", "1", "0")))
amenities2 = amenities2 %>% mutate(Gym=as.factor(ifelse(Gym=="YES", "1", "0")))
```

```
amenities2 = amenities2 %>% mutate(Tennis=as.factor(ifelse(Tennis=="YES", "1", "0")))
amenities2 = amenities2 %>% mutate(Spa=as.factor(ifelse(Spa=="YES", "1", "0")))
amenities2 = amenities2 %>% mutate(Casino=as.factor(ifelse(Casino=="YES", "1", "0")))
amenities2 = amenities2 %>% mutate(Internet=as.factor(ifelse(Internet=="YES", "1", "0")))
```

Linear Regression: Lets see how influential the 6 amenities are on the hotel's score. To do this, lets regress Pool, Gym, Tennis, Spa, Casino, and Internet on Score. We will call our linear model 'am\_model'

```
#am_model1<-glm(Score~Pool+Gym+Tennis+Spa+Casino+Internet, data=amenities1,
family='binomial') # run regression
#summary(am_model1) # summarize linear model

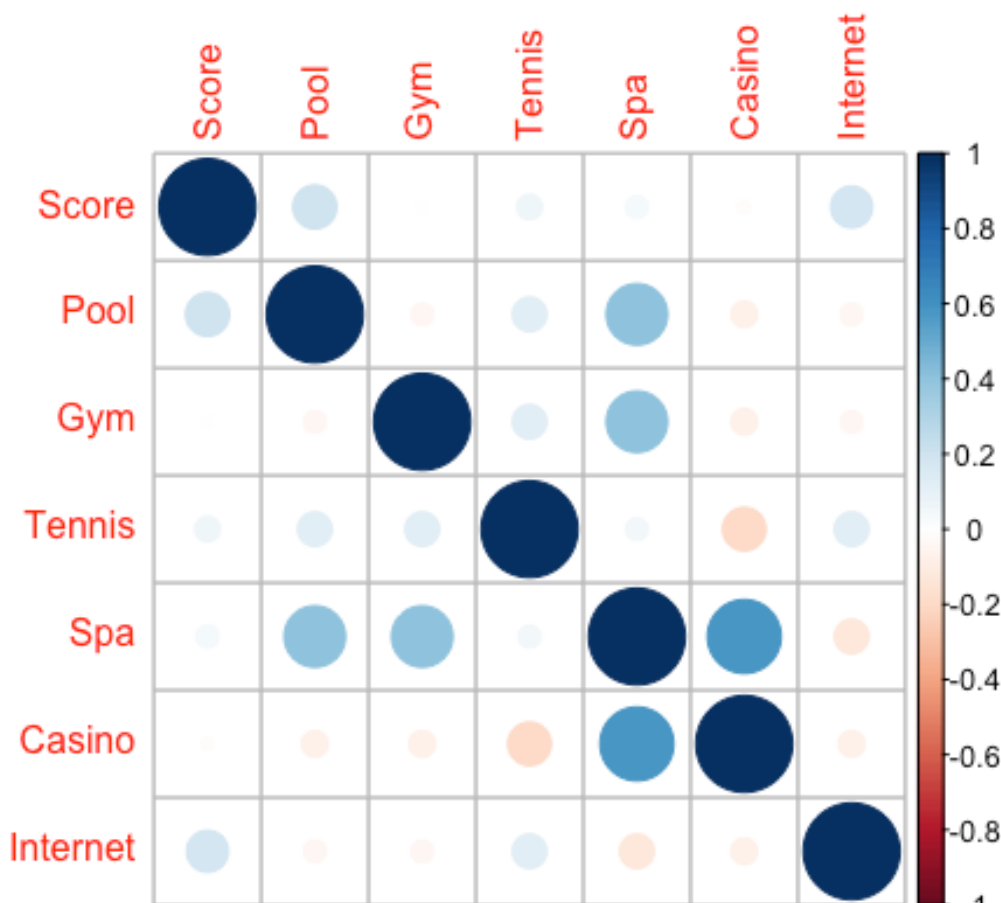
#am_model2<-glm(Score~., data=amenities2, family='binomial') # run regression
#summary(am_model2) # summarize linear model
```

The first argument in the glm (generalized linear model) function is the formula. You place the independent variable on the left of the tilde and the regressors or dependent variables on the right. The second argument simply specifies that we are using our amenities data set and the third specifies that our data is binomial (ie uses logit link function).

Look at the summary table and see if you can interpret the different columns. Recall: \* Estimate (Coefficient): \* Std. Error: \* T Value: \* P(>|t|)

Correlation Matrix: Are the variables related to each other in any way? What variables are most related to Score?

```
#amenities1[,3:9]
cor.amen<-cor(amenities1[,3:9], use="complete", method="pearson")
corrplot(cor.amen)
```



TRY:

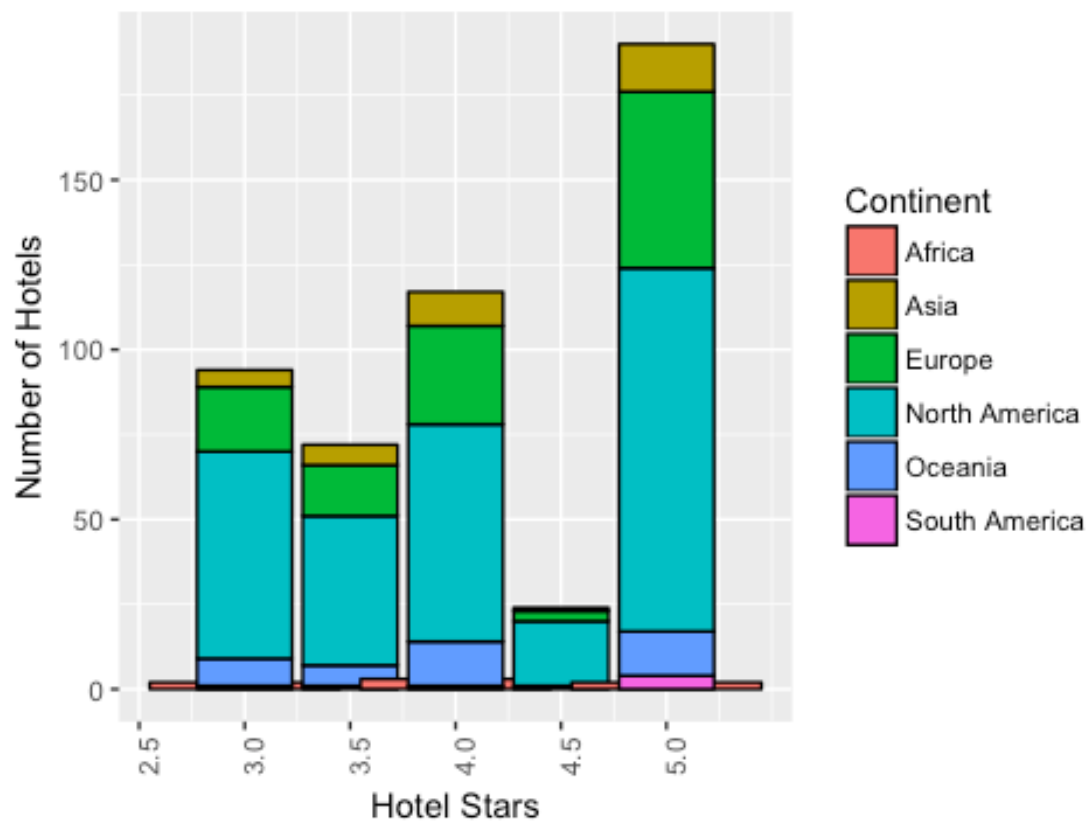
`geom_bar(stat='identity')???`

Lets make a graph looking at how many hotels f

```
g<- ggplot(df)+
  geom_bar(aes(x=Stars, fill = Continent), color="black") + #create bars & fill color
  according to continent
  theme(plot.title=element_text(hjust=0.5), # center title
        axis.text.x=element_text(angle=90, hjust=1)) + # tilt the x-axis labels 90
  degrees to the left
  xlab("Hotel Stars") + # add Label x axis
  ylab("Number of Hotels") + # add Label y axis
  scale_x_continuous(breaks= seq(2,5,by=0.5), # make our tick marks count half stars
                    labels = c('2.0', '2.5', '3.0', '3.5','4.0', '4.5', '5.0')) +
  ggtitle("Hotels Per Rating and Continent") #title ggplot
g #shows plot

## Warning: position_stack requires non-overlapping x intervals
```

Hotels Per Rating and Continent



`#ggplotly(g) # turn into plotly plot (interactive)`